

Yinbo Wang

School of Statistics, Renmin University of China | Beijing, China

audiofool@ruc.edu.cn | audiofool.blog | [GitHub: Audiofool934](https://github.com/Audiofool934)

Education

Renmin University of China

2023–Present

Bachelor of Science in Statistics (Statistics and Data Science Honors Program)

GPA: 3.71/4.0 | Weighted average: 88.24/100 | Expected graduation: June 2027

Research Interests

Multimodal generative modeling for music, audio, and video; audio-visual learning and controllable generation

Publications and Manuscripts

Xihua Wang*, Yinbo Wang*, Jingchao Zhang*, and Ruihua Song. **Meric: A Unified Framework for Multimodal Music Generation and Retrieval via Representation Space Anchoring.**

Accepted to *ECCV 2026*. *Equal contribution. [\[Project\]](#)

Xin Cheng, Xihua Wang, Ying Ba, Yuyue Wang, Kaisi Guan, Yinbo Wang, Wenpu Li, and Ruihua Song. **SyncDPO: Enhancing Temporal Synchronization in Video-Audio Joint Generation via Preference Learning.** Preprint, 2026; submitted to *NeurIPS 2026*. [\[Paper\]](#) [\[Project\]](#)

Research Experience

Research Assistant, AIMind

Jul 2025–Present

Renmin University of China | Advisor: Prof. Ruihua Song

- Led the development of **Meric**, a unified framework for image- and video-conditioned music generation and cross-modal retrieval through representation-space anchoring.
- Built and curated multimodal training data, designed model components, and conducted end-to-end flow-matching training and evaluation.
- Contributed dataset construction and curation to a project on synchronized audio-visual video generation.

Selected Projects

CC-Copilot. Open-source supervision tool for coding agents, with cross-session monitoring, citation-backed summaries, and grounded conversational queries

Timbreloom. Local speech and video-dubbing workbench integrating ASR, voice cloning, voice conversion, and resumable audio generation

Deep Blue. Research-agent harness for bounded experiment execution, protocol-bound evaluation, and auditable evidence tracking

Moondog. Personal music agent with listening-history memory, user-correctable preference profiles, and grounded playlist planning

Academic Activities

X-AGI Conference. Organizer and website designer (2026 edition upcoming).

2025; 2026

GSAI-ML. Attended one semester of group meetings in Prof. Chongxuan Li's group, covering diffusion models, video generation, and reinforcement learning.

Nov 2024–Jul 2025

GeWu-Lab. Explored audio-visual editing.

Jul 2024–Oct 2024

Selected Coursework

Mathematics and Statistics: Mathematical Analysis, Advanced Algebra, Real Analysis, Probability Theory, Mathematical Statistics, Optimization, Regression Analysis, Stochastic Processes, and Statistical Computing

CS and Machine Learning: Data Structures and Algorithms, Parallel Computing, Digital Signal Processing, Machine Learning, Deep Learning, Reinforcement Learning, and Probabilistic Graphical Models

Honors and Languages

Honors: Excellence in Study Scholarship, Third Class (2024); Beijing Merit Student (San Hao)

English: TOEFL iBT: 5.0/6.0 (100/120 comparable; CEFR C1); February 7, 2026